

Dirichlet Mixtures, the Dirichlet Process, and the Structure of Protein Space

Viet-An Nguyen¹, Jordan Boyd-Graber² and Stephen F. Altschul^{3,*}

¹ Department of Computer Science and UMIACS,
University of Maryland
College Park, MD 20740, United States
`vietan@cs.umd.edu`

² iSchool and UMIACS,
University of Maryland
College Park, MD 20740, United States
`jbg@umiacs.umd.edu`

³ National Center for Biotechnology Information,
National Library of Medicine, National Institutes of Health,
Bethesda, MD 20894, United States
Tel.: +1 (301) 435-7803
Fax: +1 (301) 480-2288
`altschul@ncbi.nlm.nih.gov`

September 29, 2012

*Corresponding author

Abstract

The Dirichlet process is used to model probability distributions that are mixtures of an unknown number of components. Amino acid frequencies at homologous positions within related proteins have been fruitfully modeled by Dirichlet mixtures, and we use the Dirichlet process to derive such mixtures with an unbounded number of components. This application of the method requires several technical innovations to sample an unbounded number of Dirichlet mixture components. The resulting Dirichlet mixtures model multiple alignment data substantially better than do previously derived ones. They consist of over 500 components, in contrast to fewer than 40 previously, and provide a novel perspective on the structure of proteins. Individual protein positions should be seen not as falling into one of several categories, but rather as arrayed near probability ridges winding through amino-acid multinomial space.

1 Introduction

Given a multiple alignment of sequences from a particular protein family, how may one estimate the amino acid frequencies found in related sequences at a specific alignment position, and thereby construct scores for adding a new sequence to the alignment? An elegant Bayesian approach to this problem was proposed in the 1990s by researchers at UCSC (Brown et al., 1993; Sjölander et al., 1996). In brief, one may model a particular position in a particular protein family by an unknown set of twenty amino acid probabilities, a point in the multinomial space Ω_{20} . Given a prior probability density P over Ω_{20} , Bayes' theorem implies a posterior density P' after the observation of several amino acids at the position in question. An estimate \vec{q} for the amino acid frequencies of the protein family at this position may then be derived by integrating P' over Ω_{20} . Although the prior density P may be of arbitrary form, it is mathematically convenient if P is assumed to be a Dirichlet distribution, or a mixture of M Dirichlet distributions.

As the number of observed amino acids at a position grows, \vec{q} converges to the observed frequencies, no matter what the prior P is. However, given a small number of observations, \vec{q} will in general be a better estimate of the actual probabilities at the protein position if the prior P accurately describes the density over Ω_{20} characteristic of real protein families. Discovering such a P given data is a problem of posterior inference. One starts with a large “gold standard” dataset S of protein multiple alignments, which are assumed to be accurate. Each “column” from these multiple alignments represents a particular position within a particular protein family, and it is really these columns that may be considered as constituting the dataset. One then seeks the maximum-likelihood Dirichlet mixture (DM).

One immediate problem arises. The likelihood of S may in general be improved by increasing the number of components of a DM until it roughly equals the number of columns in S . Doing so, however, leads to the classic problem of overfitting the data, which causes degraded predictions on new data. One solution is to apply the Minimum Description Length principle (Grünwald, 2007), and seek instead to minimize the “total description length” $\text{COMP}(\mathfrak{D}_M) + \text{DL}(S|\theta)$ (Ye et al., 2011b).^a The first term of this expression is the “complexity” of the model \mathfrak{D}_M consisting of all M -component DMs; this can be understood as the log of the effective number of independent theories \mathfrak{D}_M contains (Grünwald, 2007). The second term is the negative log likelihood of S implied by the maximum-likelihood θ drawn from \mathfrak{D}_M . Although no feasible algorithm for minimizing $\text{DL}(S|\theta)$ is known, approximations may be found using approaches based on expectation maximization (Brown et al., 1993; Sjölander et al., 1996) or Gibbs sampling (Ye et al., 2011b).

An alternative approach which never fixes M , but treats the number of components as unknown, is possible using nonparametric Bayesian models. One such model is the Dirichlet process (DP), which we apply to multiple alignment data. In brief, the DP allows us to create a generalized prior probability density over the space of DMs with an unlimited number of components. Posterior inference using a Gibbs sampling algorithm moves naturally among mixture models with varying numbers of components. The DP and its generalization, the Pitman-Yor distribution (Pitman and Yor, 1997), have been applied previously to Gaussian mixtures (Antoniak, 1974), mixtures of multinomials (Hardisty et al., 2010), admixtures of multinomials (Teh et al., 2006), time-dependent mixtures of multinomials (Beal et al., 2002), and mixtures of linear models (Hannah et al., 2011), but not to Dirichlet mixtures. In describing probability

^aOther validation approaches that are robust to overfitting include held-out perplexity (Blei et al., 2003) or extrinsic evaluation.

densities over Ω_L , Dirichlet mixtures have much greater flexibility than do multinomial mixtures. An individual multinomial component can model only probability concentrated at a specific location in Ω_L , whereas a single Dirichlet component can model densities that are arbitrarily concentrated around such a location, and even densities with most of their mass near the boundaries of Ω_L . The components of a Dirichlet mixture may have probability densities of variable concentration. Thus, for example, one component can favor positions with a fairly precise amino acid probability signature, whereas another can favor positions that contain hydrophobic amino acids, but only one or a small subset of them.

When used to analyze the same dataset for which a previous study (Ye *et al.*, 2011b) yielded a 35-component DM, our DP-based Gibbs sampling algorithm yields substantially improved solutions with over 500 components. Such large DMs may be cumbersome for practical algorithms, but a specified trade-off between component number and total description length can be used to select a DM with fewer components.

Of perhaps greater interest is the perspective on the structure of protein space provided by DMs with many components. The DM formalism suggests, at first, the metaphor of a small number of probability hills in Ω_{20} , corresponding to different types of protein positions—hydrophobic, aromatic, charged, etc. However, the density implied by the many-component DMs we derive is dominated by a continuous probability ridge winding through Ω_{20} . This may provide a new perspective on how selective pressures are felt at individual protein positions.

2 Methods

We here describe the mathematical underpinnings of our approach, providing a brief review of standard material, and devoting more detailed discussion to less familiar or novel methods.

2.1 Multinomial space

A multinomial probability distribution on an alphabet with L letters is a vector with L positive components that sum to 1, and the space of all possible multinomials is the simplex Ω_L . Due to the constraints on the vector components, Ω_L is finite and has $(L - 1)$ degrees of freedom. For example, Ω_3 is the 2-dimensional equilateral triangle, embedded in Euclidean 3-space, with vertices $(1,0,0)$, $(0,1,0)$ and $(0,0,1)$. We will be interested primarily in the standard amino acid alphabet, and therefore in the 19-dimensional space Ω_{20} .

2.2 The Dirichlet distribution

For an alphabet of L letters, a Dirichlet distribution D is a probability density over Ω_L , parameterized by an L -dimensional vector $\vec{\alpha}$ of positive real numbers greater than zero; it is convenient to define α as $\sum_{j=1}^L \alpha_j$. The density of D at \vec{x} is given by

$$D(\vec{x} \mid \vec{\alpha}) \equiv Z \prod_{j=1}^L x_j^{\alpha_j - 1}, \quad (1)$$

where the normalizing scalar $Z \equiv \Gamma(\alpha) / \prod_{j=1}^L \Gamma(\alpha_j)$ is chosen so that integrating D over Ω_L yields 1. The expectation or mean of \vec{x} under the density of equation (1) is the multinomial distribution parameterized by $\vec{q} \equiv \vec{\alpha} / \alpha$. It is frequently useful to

write D 's parameters in the form $\alpha\vec{q}$, with $\alpha \in (0, \infty)$, and $\vec{q} \in \Omega_L$. When we use this alternative parametrization for D , we write it as (\vec{q}, α) . Intuitively, one may visualize a Dirichlet distribution as a probability hill in Ω_L , centered at \vec{q} , and with greater α corresponding to greater concentration of probability mass near \vec{q} . For α near 0, the “hill” in fact becomes a trough, with most probability concentrated near the boundaries of Ω_L . Thus, such Dirichlet distributions favor sparse multinomials where only a few letters have non-negligible probability (Figure 1).

Figure 1

2.3 Mixture models

Given a proposed set of observations, a theory may be thought of as assigning probabilities to all possible datasets or outcomes. If a theory has a particular set of adjustable parameters, we call the set of all such theories a model. More generally, we may wish to consider multiple models, usually nested, characterized by different numbers or sets of parameters.

Mixture models are a formalism frequently used to discover clustering patterns in data. In a mixture model, all observations are associated with clusters, each of which has a corresponding probabilistic mixture “component” that explains its data. For example, multinomial mixture models are frequently used in text analysis (Lewis, 1998).

Multinomial mixtures have difficulty modeling many probability densities over Ω_L , because the $L - 1$ free parameters of an individual multinomial component can only describe probability concentrated at a specific location \vec{q} . In contrast, with the addition of the single extra parameter α , a Dirichlet component can describe probability densities of arbitrary concentration around \vec{q} , including, when α is small, densities that favor sparse multinomials (Figure 1). This greatly enhanced flexibility allows a Dirich-

let mixture to model most real-world probability densities over Ω_L much better than can a multinomial mixture with many times as many components.

An M -component DM is a probability density over Ω_L , defined as the weighted sum of M Dirichlet distributions, called Dirichlet components. Such a mixture has $ML + M - 1$ free parameters. Each of the M Dirichlet components contributes L “Dirichlet parameters”. In addition, the weights \vec{w} or “mixture parameters” are M positive real numbers that sum to 1, only $M - 1$ of which are independent. A DM may be thought of as a superposition of M probability hills in Ω_L , each with its particular volume, center of mass and concentration.

2.4 The Dirichlet process

When seeking a theory for a set of data, a difficulty is that theories with more parameters generally can explain the data better, but overfitting can result in poor predictions on future data. One approach to this problem is the Minimum Description Length principle (Grünwald, 2007), which explicitly favors theories drawn from mixture models with fewer components (Ye *et al.*, 2011b). An alternative approach is provided by the Dirichlet process (DP), which effectively subsumes in a single formalism mixture models with an arbitrary number of components. A mathematically detailed description of the Dirichlet process (DP) can be found elsewhere (Antoniak, 1974; Pitman and Yor, 1997; Müller and Quintana, 2004); here we will review only its essentials.

The DP generally is applied to problems where data are postulated to be well-modeled as generated by a mixture of multiple instances (often called “atoms” but here called “components”) of an underlying distribution of known parametric form. In the DP formalism, every mixture consists of a countably infinite number of components, each with its own weight and set of component parameters. In essence, a DP defines a

generalized probability distribution over this infinite-dimensional space of mixtures.

Two elements completely specify a DP:

1. A “base” probability distribution H over the space of component parameters.
For example, if the components are Gaussians on \mathbb{R} with unit variance, H is a specified distribution for their means.
2. A positive real parameter, which implicitly defines a probability distribution on component weights. This parameter is usually called α , but we will call it γ here to avoid the potential confusion arising from the multiple distinct uses we make of Dirichlet distributions. As we will see, the smaller γ , the greater the implied concentration of weight in a few components.

Analysis using the DP is Bayesian. A DP is used to define a prior over mixture distributions which, when combined with observed data, implies a posterior for the weights and component parameters of these mixtures. A special feature of this inference is that, although all mixtures are assumed to have a countably infinite number of components, only a finite number can ever explain a given set of data. The posterior distribution thus differs from the prior only for finitely many components. Bayesian analysis allows one to estimate the number of these components, as well as their associated weights and component parameters.

2.5 The Chinese restaurant process

The “Chinese restaurant process” (CRP) (Ferguson, 1973) is closely related to the Dirichlet process and is useful for understanding the properties of the DP, as well as for posterior inference. The metaphor in the name refers to a restaurant with an unbounded number of tables.

The Chinese restaurant is patronized by customers. Each customer represents an i.i.d. draw from a distribution G drawn from a Dirichlet process $DP(\gamma, H)$.^b Each customer sits at one of the tables, and when customers sit at the same table it means they are associated with the same component, drawn from the base distribution H . In our application, a customer represents a multiple-alignment column, and a table represents a DM component, (\vec{q}_k, α_k) .

Draws from the Dirichlet process are exchangeable (Aldous, 1985), so each customer can be viewed as the “last” customer to enter the restaurant. When customers enter, they choose to sit either at a new table or at one that is already occupied. This choice is made randomly, but with each occupied table selected with probability proportional to the number of people already seated there, and a new table selected with probability proportional to the parameter γ . It is evident that smaller values for γ imply a greater concentration of customers at a small number of tables.

The exchangeability of the Dirichlet process is important for Gibbs sampling inference, because it allows us to condition one column’s component assignment on the other columns’ assignments.

2.6 A base distribution for Dirichlet-component parameters

Specifying a DP for Dirichlet mixtures requires specifying a base distribution H over the space of Dirichlet parameters. Rather than defining H on the standard Dirichlet parameters $\vec{\alpha} \in \mathbb{R}^{+L}$, we find it more natural to define it on the alternative parameters (\vec{q}, α) . Specifically, we propose $H \equiv (H_1, H_2)$, where H_1 and H_2 are independent distributions for \vec{q} and α .

^bNote that the Chinese restaurant does not model a specific measure G ; it only models draws from G and integrates over all possible G . However, this representation is sufficient for our purposes. For a constructive definition of the Dirichlet process, see (Sethuraman, 1994).

Because $\vec{q} \in \Omega_{20}$, a natural base distribution H_1 for \vec{q} is itself Dirichlet. Furthermore, because we seek a DM that describes protein columns, it is appropriate to choose H_1 's center of mass to be \vec{p} , the “background” amino acid frequencies typical for proteins. This leaves only the single concentration parameter, which we will call β . In short, we propose choosing H_1 to be the Dirichlet distribution with parameters (\vec{p}, β) .

When specifying H_2 , the base distribution for $\alpha \in (0, \infty)$, we will see that it is convenient if we require H_2 to have a long, uninformative tail. An exponential function of the form $H_2 \equiv \lambda e^{-\lambda\alpha}$, with λ small, serves the purpose, and the precise value of λ will be irrelevant.

By choosing \vec{p} as the center of mass for H_1 , and requiring H_2 to have a long tail, the base distribution $H \equiv (H_1, H_2)$ we propose for Dirichlet-component parameters has, in effect, only the one free parameter β , as results are insensitive to the choice of λ . This, in conjunction with the parameter γ , completes our specification of a DP for Dirichlet mixtures. We will discuss in the Results section the effects of different choices for β and γ .

2.7 Model

Figure 2

To review, we posit the following generative process for observed data:

- We draw component $k \in [1, \infty)$ of the Dirichlet process from the base distribution; this draw has two parts:
 - the component's mean \vec{q}_k is sampled from $\text{Dirichlet}(\vec{p}, \beta)$;
 - the component's concentration α_k is sampled from $\text{Exponential}(\lambda)$, which is equivalent to a Gamma distribution with shape = 1 and rate = λ .

- We draw weights \vec{w} for all of the Dirichlet process components from $\text{GEM}(\gamma)$.^c
- For each column $i \in [1, n]$:
 - We draw a component assignment z_i from the distribution \vec{w} ;
 - We draw a multinomial distribution $\vec{\phi}_i$ over letters from $\text{Dirichlet}(\vec{q}_{z_i}, \alpha_{z_i})$;
 - We draw the letters of column i from $\text{Multinomial}(\vec{\phi}_i)$, resulting in the observation vector \vec{x}_i , with associated letter count vector \vec{c}_i .

We assume that this process created the observed columns and use posterior inference, described in the rest of the section, to uncover the latent variables that best explain the observed data. The generative process may be expressed using the graphical model in Figure 2.

At this point, we pause to recognize that our terminology has potential confusions. Our model has three different uses of the word “Dirichlet”. One is a Dirichlet process, and two are ordinary Dirichlet distributions:

- At the top level, a Dirichlet process gives us a countably infinite number of components. This is a nonparametric Bayesian distribution over distributions.
- Each component k of the Dirichlet process is itself a Dirichlet distribution, parameterized by (\vec{q}_k, α_k) . This is a finite distribution over Ω_{20} . Columns are generated from a multinomial drawn from this distribution.
- The mean of each component’s Dirichlet distribution is itself drawn from a Dirichlet distribution parameterized by (\vec{p}, β) .

^cThe vector \vec{w} is a point on the infinite simplex (Sethuraman, 1994) and GEM stands for Griffiths (Griffiths, 1980), Engen (Engen, 1975), and McCloskey (McCloskey, 1965). This, along with the separate component draws, form a constructive “stick breaking” definition of the Dirichlet process.

2.8 MCMC inference for Dirichlet process Dirichlet mixtures

The Gibbs sampling algorithm for DMs described in (Ye et al., 2011b) assumed a fixed number of components M . The algorithm alternated between a first stage, in which the component assignments z_i for columns were chosen by Gibbs sampling, conditioned on the complete set of component parameters, and a second stage, in which each component’s parameters were updated, using maximum-likelihood estimates based on the columns associated with that component. Our approach here, although similar in many ways, has a few key differences.

Like the previous approach, our Gibbs sampler forms a Markov chain over assignments to components $\mathbf{Z} \equiv \{z_1, z_2, \dots, z_n\}$ and component parameters $\{(\vec{q}_1, \alpha_1), (\vec{q}_2, \alpha_2) \dots\}$. However, unlike the previous approach, the number of components is not fixed. Components can be both lost and created, but only finitely many components are ever used to describe a given dataset. Specifically, before being assigned stochastically to a new component, by means of the latent variable z_i , a column i is first removed from an existing one, and if this component is left with no associated columns, it is abolished. Then, for the column’s new assignment, it may choose among the existing components, but it may also start a new one. Unlike in (Ye et al., 2011b), this sampling is conditioned on the current component assignments of all other columns, rather than on those assignments only from the previous round. During the algorithm’s second stage, we sample component parameters, rather than update them by maximum-likelihood estimation. There are various ways in which one may initialize the algorithm, but the simple expedient of assigning all columns to a single component at the start does not appear to cause any difficulties. We describe in greater detail below various technical aspects of these modifications to the algorithm of (Ye et al., 2011b).

2.9 Sampling an existing or new component for a column

Our DP-sampling algorithm creates a Markov chain of component assignments and component parameters. While sampling component assignments, we assume that the Dirichlet parameters $\vec{\alpha}_k$ associated with an existing component k remain fixed. However, the number of columns n_k associated with component k may change.

Gibbs sampling conditions a column’s assignment z_i on the other columns’ assignments and the component parameters. This is where the exchangeability of the Chinese restaurant process (Section 2.5) is advantageous. Computing the conditional distribution is equivalent to removing column i from its current component and then assigning it to an existing component or a completely new one. Assuming the observations \vec{x}_i of column i contain c_i total amino acids, with the amino acid counts given by the vector \vec{c}_i , the likelihood for an existing component k is proportional to

$$p(z_i = k \mid \mathbf{Z}^{-i}, \vec{q}_k, \alpha_k, \vec{x}_i) \propto n_k \cdot \frac{\Gamma(\alpha_k)}{\Gamma(\alpha_k + c_i)} \prod_{j=1}^{20} \frac{\Gamma(\alpha_k q_{k,j} + c_{i,j})}{\Gamma(\alpha_k q_{k,j})} \quad (2)$$

(Brown et al., 1993; Sjölander et al., 1996; Altschul et al., 2010; Ye et al., 2011b). Here, \mathbf{Z}^{-i} denotes the set of component assignments for all columns except column i . In the Chinese restaurant metaphor, this corresponds to sitting at an existing table.

In addition to being associated with an existing table, there is also a probability of sitting at a new table; this happens with probability proportional to γ . Because this is a new table the component parameters are unknown, however we still must calculate the probability of the column’s observations \vec{x}_i . The proper Bayesian approach is to integrate this likelihood over all possible Dirichlet distributions (\vec{q}, α) , given the base distribution $H \equiv (H_1, H_2)$, so that

$$p(z_i = \text{new} \mid \vec{x}_i, H_1, H_2, \gamma)$$

$$\begin{aligned} &\propto \gamma \cdot p(\vec{x}_i \mid H_1, H_2) \\ &= \gamma \int_0^\infty H_2(\alpha) d\alpha \int_{\Omega_{20}} H_1(\vec{q}) d\vec{q} \int_{\Omega_{20}} p(\vec{x}_i \mid \vec{\phi}) D(\vec{\phi} \mid \alpha \vec{q}) d\vec{\phi}, \end{aligned} \quad (3)$$

where $p(\vec{x}_i \mid \vec{\phi})$ is the probability of observing column i given the multinomial distribution $\vec{\phi}$.

This is where our choice of H_2 is first advantageous. Because H_2 is a function that decays very slowly in α , almost all of the mass of such a density is contributed by large α , for which the corresponding Dirichlet distributions can be considered delta functions at \vec{q} . This reduces the right side of equation (3) to

$$\gamma \int_{\Omega_{20}} p(\vec{x}_i \mid \vec{q}) H_1(\vec{q}) d\vec{q}, \quad (4)$$

which, analogously to before (Brown et al., 1993; Sjölander et al., 1996; Altschul et al., 2010; Ye et al., 2011b), is just the probability of observing the amino acid vector \vec{x}_i given the Dirichlet distribution (\vec{p}, β) over multinomials. Thus,

$$p(z_i = \text{new} \mid \vec{x}_i, H_1, H_2, \gamma) \propto \gamma \cdot \frac{\Gamma(\beta)}{\Gamma(\beta + c_i)} \prod_{j=1}^{20} \frac{\Gamma(\beta p_j + c_{i,j})}{\Gamma(\beta p_j)}. \quad (5)$$

Given Equation 2 and Equation 5, we may now sample column i into an existing or a new component. If a new component is selected, our final problem is how to assign it a set of Dirichlet parameters. Here, we simply use the sampling method described in the next section, but applied to a component with only a single associated column.

2.10 Sampling Dirichlet component parameters

In addition to sampling column assignments, we must also sample the Dirichlet component parameters $\{(\vec{q}_k, \alpha_k)\}$. As in prior work (Ye et al., 2011b), we take a coordinate-wise approach for sampling these parameters: for each component k , first sampling \vec{q}_k , and then sampling α_k .

We sample \vec{q}_k^* from a Dirichlet distribution with parameters $\beta\vec{p} + \vec{C}_k$, where \vec{C}_k is the aggregate observation vector, summing over all columns associated with component k . This approximates the true posterior of \vec{q}_k under the maximal path assumption (Walach, 2008).

Given our sampled \vec{q}_k^* , the column data yield an analytic formula for the log-likelihood function $\mathfrak{L}(\alpha_k)$ as a function of α_k , as well as for its first and second derivatives (Minka, 2000; Ye et al., 2011b). If, as suggested above, the prior on α_k takes the form $\lambda e^{-\lambda\alpha_k}$ then the posterior log-likelihood is, up to a constant, $\mathfrak{L}(\alpha_k) - \lambda\alpha_k$. Assuming λ is small permits us to ignore the second term, and to sample α_k^* with reference only to $\mathfrak{L}(\alpha_k)$. For certain special cases, \mathfrak{L} has a supremum at 0 or ∞ , and in these instances one may set α_k^* respectively to a fixed small or large number. (In either case, the likelihood subsequently implied for amino acid count vectors is insensitive to the precise value chosen for α_k^* .) Otherwise, it has been postulated but not proved (Ye et al., 2011b) that \mathfrak{L} has a unique maximum at $\hat{\alpha}_k$, which it is easy to locate using Newton’s method. If \mathfrak{L} ’s second derivative at $\hat{\alpha}_k$ is $-X$, we can use the Laplace approximation to sample α_k^* from a normal distribution with mean $\hat{\alpha}_k$ and variance $1/X$.

2.11 Refinements to inference

While the above methods were used to generate the results reported in Section 3 up to Section 3.3, we describe here alternative methods for those interested in more refined inference techniques that allow a joint search over more of the model’s parameters.

In Section 3, we employ a grid search to determine the Dirichlet process parameter γ that optimizes the MDL, an objective function previously proposed for Dirichlet mixture models used in computational biology (Ye *et al.*, 2011b). Alternative objective functions could be classification performance, interpretability, or biological plausibility.

The objective function of likelihood can be optimized with approximate inference or, more specifically, using slice sampling (Neal, 2003). Slice sampling is a general MCMC algorithm to draw random samples from an unknown probabilistic distribution by sampling uniformly from the region under the variable’s density function. In our case, we would like to sample the hyperparameter γ from $p(\gamma \mid \mathbf{A}, \mathbf{Q}, \mathbf{Z}, \mathbf{X}; \beta, \lambda, \vec{p})$ ^d. This density is proportional to the data likelihood $p(\mathbf{X} \mid \mathbf{Z}, \mathbf{Q}, \mathbf{A}; \beta, \gamma, \lambda, \vec{p})$ where

- $\mathbf{X} = \{\vec{x}_i \mid i \in [1, n]\}$ denotes all observed amino acid columns.
- $\mathbf{Z} = \{z_i \mid i \in [1, n]\}$ denotes the component assignments for all columns.
- $\mathbf{Q} = \{\vec{q}_k \mid k \in [1, K^+]\}$ denotes the centers of mass for all components.
- $\mathbf{A} = \{\alpha_k \mid k = [1, K^+]\}$ denotes the concentration parameters for all components.

Here, K^+ is the current number of components^e. This density can be rewritten as:

$$p(\mathbf{X} \mid \mathbf{Z}, \mathbf{Q}, \mathbf{A}; \beta, \gamma, \lambda, \vec{p}) = p(\mathbf{X} \mid \mathbf{Z}, \mathbf{Q}, \mathbf{A}) \cdot p(\mathbf{Z} \mid \gamma) \cdot p(\mathbf{A} \mid \lambda) \cdot p(\mathbf{Q} \mid \beta, \vec{p}). \quad (6)$$

^dSlice sampling can be applied to more general parameters. A more general alternative would be to slice sample a vector $v \equiv (\beta, \gamma, \lambda)$ (Wallach, 2008).

^eThe superscript ⁺ is to denote that the number of components is unbounded and varies during our sampling process.

The right hand side of Equation 6 consists of the following factors:

- The likelihood of tables' parameters given the observed columns:

$$p(\mathbf{X} \mid \mathbf{Z}, \mathbf{A}, \mathbf{Q}) = \prod_{i=1}^n p(\vec{x}_i \mid z_i = k, \alpha_k, \vec{q}_k) = \prod_{i=1}^n \frac{\Gamma(\alpha_k)}{\Gamma(\alpha_k + c_i)} \prod_{j=1}^L \frac{\Gamma(\alpha_k q_{k,j} + c_{i,j})}{\Gamma(\alpha_k q_{k,j})}. \quad (7)$$

- The joint distribution of table assignments^f:

$$p(\mathbf{Z} \mid \gamma) = \gamma^{K^+} \frac{\prod_{k=1}^{K^+} (n_k - 1)!}{\prod_{i=1}^n (i - 1 + \gamma)} \approx \gamma^{K^+} \frac{\prod_{k=1}^{K^+} \Gamma(n_k)}{\frac{\Gamma(n+\gamma)}{\Gamma(\gamma)}}. \quad (8)$$

Note that n_k is the number of columns currently assigned to table k and n is the total number of columns. The expression in Equation 8 is exact when γ is an integer.

- The likelihood of λ given the current values of concentration parameters \mathbf{A} :

$$p(\mathbf{A} \mid \lambda) = \prod_{k=1}^{K^+} \lambda e^{-\lambda \alpha_k}. \quad (9)$$

- The likelihood of β given the current values of centers of mass \mathbf{Q} :

$$p(\mathbf{Q} \mid \beta, \vec{p}) = \prod_{k=1}^{K^+} \frac{\Gamma(\beta)}{\prod_{j=1}^L \Gamma(\beta p_j)} \prod_{j=1}^L q_{k,j}^{\beta p_j - 1}. \quad (10)$$

The detailed pseudo-code of the slice sampling used is shown in Algorithm 1 (Walach, 2008). The algorithm requires two inputs: (1) R is the number of iterations and is typically set to 10, and (2) σ is the step size and set to one tenth the current value

^fFor detailed derivation of Equation 8, refer to (Gershman and Blei, 2012).

of γ . For notational convenience, we use $(\gamma^{left}, \gamma^{right})$ to denote the range around γ we sample from, and use $f(\gamma)$ to denote $p(\gamma \mid \mathbf{A}, \mathbf{Q}, \mathbf{Z}, \mathbf{X}; \beta, \lambda, \vec{p})$.

Algorithm 1: Pseudo-code for slice sampling algorithm used.

Input: R : number of iterations. σ : step size.

Initialize γ

foreach $r \in [1, R]$ **do**

 Draw $u' \sim \text{Uniform}(0, f(\gamma))$;

 Draw $v \sim \text{Uniform}(0, 1)$

$\gamma^{left} \leftarrow \gamma - v\sigma$

$\gamma^{right} \leftarrow \gamma^{left} + \sigma$

while **true** **do**

 Draw $\gamma' \sim \text{Uniform}(\gamma^{left}, \gamma^{right})$

if $\frac{f(\gamma')}{f(\gamma)} > u'$ **then**

\perp break;

else

if $\gamma' < \gamma$ **then**

\perp $\gamma^{left} \leftarrow \gamma'$

else

\perp $\gamma^{right} \leftarrow \gamma'$

3 Results

The research group at UCSC that first proposed Dirichlet mixtures for protein analysis currently makes a number of multiple alignment datasets available on their web site, <http://compbio.soe.ucsc.edu/dirichlets/index.html>. We consider their dataset “diverse-1216-uw”, called S_{UCSC} here, which was studied also in (Ye *et al.*, 2011b). S_{UCSC} consists of 23,903,805 amino acids arranged into 314,585 columns, and thus containing a mean of approximately 76.0 amino acids per column.

3.1 The quality of Dirichlet mixtures

It is useful to have an objective measure for the quality of a DM, and for this purpose we turn to the Minimum Description Length (MDL) principle (Grünwald, 2007), whose essentials we review here.

One may define the description length of a dataset S , given a theory θ , as $\text{DL}(S|\theta) \equiv -\log_2 P_\theta(S)$, i.e. the negative log of the probability for the dataset implied by the theory. Because the logarithm is to the base 2, DL is said to be expressed in bits. This definition may be extended to a model \mathfrak{M} by defining $\text{DL}(S|\mathfrak{M}) \equiv \inf_{\theta \in \mathfrak{M}} \text{DL}(S|\theta)$.

If one wishes to find the model that best describes a set of data, using DL alone as a criterion is problematic because, for nested models, increasing the number of parameters can only decrease DL. Accordingly, MDL theory introduces the formal concept of the complexity of a model $\text{COMP}(\mathfrak{M})$ (Grünwald, 2007), which may be thought of, intuitively, as the log of the effective number of independent theories the model contains. The MDL principle then asserts that the model best justified by a set of data is that which minimizes $\text{COMP}(\mathfrak{M}) + \text{DL}(S|\mathfrak{M})$. In essence, the principle supports a theory drawn from a model of greater complexity only when this increased

complexity is offset by a sufficient decrease in data description length.

To select among Dirichlet mixture models \mathfrak{D}_M with a variable number M of components using the MDL principle, one must be able at least to approximate both $\text{COMP}(\mathfrak{D}_M)$ and $\text{DL}(S|\mathfrak{D}_M)$. Heuristic arguments (Ye *et al.*, 2011b) have extended to $\text{COMP}(\mathfrak{D}_M)$ an analytic formula for the complexity of a single-component Dirichlet model (Yu and Altschul, 2011). Calculating $\text{DL}(S|\mathfrak{D}_M)$ entails finding the maximum-likelihood M -component DM. This is an instance of the classic hard problem of optimization within a rough but correlated high-dimensional space, and approximation algorithms have been based on expectation maximization (EM) (Brown *et al.*, 1993) and Gibbs sampling (Ye *et al.*, 2011b).

For the sake of analysis, we may treat our DP-sampler as simply an improved algorithm for finding DMs that minimize total description length. To evaluate a particular DM, we compare it to the baseline multinomial model in which all amino acids are drawn randomly according to background probabilities \vec{p} inferred from the data. For this model, the description length of S_{UCSC} is 99,604,971 bits, and the complexity of the model is 206 bits (Ye *et al.*, 2011b), so the total description length can be expressed as 4.1669 bits/amino acid. We assess a DM by the decrease Δ (bits/a.a.) in total description length it implies with respect to this baseline, and use Δ as an objective function of mixture quality.

3.2 The optimal number of Dirichlet components

Our DP-sampling procedure does not converge on a unique DM, and selecting different parameters β and γ of course yields different results. However, given a set of real protein multiple alignment data, the results produced by the procedure after several hundred iterations share various broad qualitative and quantitative features, which we

describe here.

At a given iteration, the current DM generated by the DP-sampler typically contains many components to which only a small number of columns are assigned. These components are particularly unstable, and with further iterations tend either to evaporate or to grow in the number of associated columns. In general, they are unsupported by the MDL principle when seeking a DM that maximizes Δ . Thus, after any given iteration, we first arrange the sampler-generated Dirichlet components in decreasing order of their number of associated columns, and then calculate the Δ implied by the DMs consisting of increasing numbers of these components. Although this greedy method does not necessarily identify the optimal subset, it provides a reasonable approximation. Typically, the MDL principle excludes sets of components to which, in aggregate, less than 2% of the columns are associated, with no single excluded component representing more than 0.05% of the columns.

Figure 3

Using a range of settings for the DP parameters β and γ described in the next section, we ran the DP-sampler on S_{UCSC} for 1000 iterations and estimated an optimal Δ after every 10th iteration. Although the best Δ frequently continued to improve past the 900th iteration, its rate of improvement always flattened out much earlier. For example, using parameters $\beta = 400$ and $\gamma = 100$, we graph in Figure 3 the best Δ found at each iteration, and in Figure 4 the number of components in the associated DM, as well as the total number of components returned. The optimal Δ of 1.0763 bits/a.a., for a 623-component DM found at iteration 940, substantially exceeds the 1.0654 bits/a.a. of a 35-component DM achieved by (Ye *et al.*, 2011b), as well as the 1.0594 bits/a.a. of the 20-component DM “dist.20comp”, derived from S_{UCSC} , that is reported on the UCSC web site. We will consider below why the DP-sampler returns DMs with so many more components than those found by earlier methods, as well as

what the size of these DMs imply about the structure of protein space.

Figure 4

3.3 Dependence on Dirichlet-process parameters

The results we obtain depend on the Dirichlet-process parameters β and γ , but their most important qualitative and quantitative features are not very sensitive to these parameters. The Dirichlet process is a Bayesian prior and its particular parameters should thus be outweighed by sufficient data. Nevertheless, it is instructive to consider the practical effects of these parameters.

Table 1

We ran the DP-sampler as described above using values for β ranging from 100 to 1000, and for γ from 5 to 100, with the results summarized in Table 1. We obtained almost equally high Δ s for β from 200 to 800, with $\gamma \geq 60$.

The β parameter specifies the concentration of the prior for the \vec{q} parameters corresponding to each Dirichlet component. When inferring \vec{q} for a particular component, the aggregate amino acid counts from its associated columns are added to β pseudo-counts. The columns in S_{UCSC} have a mean of 76 amino acids, so the seemingly large value of $\beta = 400$ in fact corresponds to only about 5 average-composition columns. This is not a very great number when, on average, > 500 columns from S_{UCSC} are assigned to each component. Because larger values for β render components with few associated columns less distinctive, they favor mixtures with fewer components, as seen in Table 1.

For all the β we tested, the best Δ found initially grew with increasing γ , but plateaued by $\gamma = 60$. Although larger values for γ favor DMs with more components, the number of components comprising the optimal results found was not very sensitive to the choice of γ . As discussed in the next section, one may avoid specifying a particular γ .

One may prefer DMs with fewer components for algorithmic reasons. In this case, it may be advantageous to use both large β and small γ ; this tends to favor DMs with fewer components, and thus to improve their corresponding Δ s. We consider below the tradeoff of Dirichlet mixture size and accuracy.

3.4 Slice sampling γ

To avoid the arbitrariness of specifying a particular value for a DP parameter, or the time involved in testing multiple values, we may use the slice sampling procedure described above. In brief, after a given iteration of the DP-sampling algorithm, we sample a new value γ' within a range centered on the current value γ . We then compute the likelihood of the current mixture model with this new γ' . If this likelihood is greater than the likelihood with the current γ , there is a high probability that we will accept this new γ' and use it in the next iteration.

We implemented this sampling procedure for the DP parameter γ , using an initial value of $\gamma = 50$, and a “burn-in” period of 25 iterations before γ is allowed to vary. For β ranging from 200 to 1000, we ran this refined algorithm for 1000 iterations; for $\beta = 100$ the program terminated after 323 iterations because the number of components it generated exceeded a limit imposed by memory constraints. In Table 2, we report for each β the mean and standard deviation for γ during the program’s last 100 iterations. We also report the optimal Δ found, its corresponding number of components, the iteration yielding this Δ , and the value of γ during this iteration.

Table 2

As can be seen, slice sampling converges on a relatively small range of values for γ , and the best Δ found is always within 0.0001 bits/a.a. of the best yielded by the multiple searches shown in Table 1, which employ fixed, specified γ . As before, $\beta \approx 400$ appears optimal, but this conclusion is now reached by a one-parameter rather than a

two-parameter search.

One may employ slice sampling to determine β as well as γ , but doing so is problematic. Although we have been using Δ as an objective function for Dirichlet mixtures, the DP-sampler is ignorant of this function, instead sampling mixtures according to their posterior likelihood, given the prior imposed by the Dirichlet process. Indeed, the DP-sampler returns mixtures with many more components than supported by the MDL principle, as seen in Figure 4. We have been able to elide this inconsistency because, for fixed β , greater posterior likelihoods for the larger mixtures correlate well with greater values for Δ . This correlation is broken, however, once β may vary. The generalized DP-sampler then prefers small β , yielding mixtures with many components, which are penalized by the model-complexity term of the MDL principle.

3.5 Tradeoff of Δ and the number of Dirichlet components

So far we have been concerned only with maximizing Δ . However, DMs are derived for use in profile-sequence, profile-profile, or multiple alignment programs (Brown *et al.*, 1993; Edgar and Sjölander, 2004; Altschul *et al.*, 2010; Ye *et al.*, 2011a), and in these applications DMs with fewer components have a speed advantage. As seen in Figure 3 and Figure 4, DMs with only slightly sub-optimal Δ can have significantly fewer components, and such DMs may well be preferred in certain circumstances.

To study this tradeoff explicitly, we recorded over all iterations of the run described in Figures 3 and 4, as well as the greedy DM-construction procedure described above, the greatest Δ found for DMs with varying numbers of components; the results are shown in Figure 5. A particular application, implementation, and preference for speed vrs. DM accuracy (i.e. Δ) can be used with such a curve to derive an optimal DM size from a software engineering perspective.

Figure 5

3.6 The topography of protein space

What do the hundreds of Dirichlet components returned by the DP-sampler imply about proteins? To study this question, it is useful to develop a representation of DMs that is easier to comprehend than would be a mere tabulation of thousands of parameters. The approach we take is to represent each component of a DM by a single line of text. On this line, we focus primarily on the component’s center-of-mass vector \vec{q} , which we represent by a string $\vec{\sigma}$ of twenty symbols, although we also report the component’s mixture parameter w and concentration parameter α numerically.

In constructing the $\vec{\sigma}$ to represent a Dirichlet component, it is useful first to order the amino acids in a manner that corresponds to their mutual similarities, even though any linear arrangement must elide some of these multi-dimensional relationships. Various orders have previously been proposed (Swanson, 1984; Brown *et al.*, 1993), but our data suggest the order “RKQEDNHWFYMLIVCTSAGP”, using the one-letter amino acid code.

Because within proteins the amino acids occur with widely differing background frequencies p_j , it is fruitful to represent a Dirichlet component’s mean “target frequencies” q_j in relation to the corresponding p_j . Accordingly, we base the symbol σ_j on the implied “log-odds score” $s_j = \log_2(q_j/p_j)$ according to the following system:

$s_j > 2$	$\sigma_j =$ The amino acid's one-letter code, in upper case
$2 \geq s_j > 1$	$\sigma_j =$ The amino acid's one-letter code, in lower case
$1 \geq s_j > 0.5$	$\sigma_j =$ “+”
$0.5 \geq s_j > -1$	$\sigma_j =$ “ ”
$-1 \geq s_j > -2$	$\sigma_j =$ “.”
$-2 \geq s_j > -4$	$\sigma_j =$ “_”
$-4 \geq s_j$	$\sigma_j =$ “=”

In other words, for a particular component, an upper case letter implies that the frequency of the corresponding amino acid is enriched vis-a-vis background by a factor greater than 4.0, while the symbol “=” means it is decreased by a factor of at least 16. We choose such a positive/negative asymmetry among the categories for defining the σ_j because $q_j \in (0, 1)$ implies an upper bound on s_j , but no lower bound.

As seen in Table 1, the DMs with greatest Δ can have over 600 components. Although we could analyze mixtures of this size, most of their important qualitative features are apparent in mixtures with many fewer components, so we will consider such a smaller DM here. As discussed above, Δ for DMs with fewer components tends to be optimized using relatively large β and small γ . Choosing $\beta = 1000$ and $\gamma = 10$, and requiring an improvement of at least 4×10^{-5} bits/a.a. in Δ for each additional component, our best result was a 134-component DM with $\Delta = 1.0732$ bits/a.a, which we call θ_{134} . The parameters associated with all components of θ_{134} are presented in Tables 3–5.

Table 3

A DM's components may be listed in arbitrary order. One reasonable choice is by decreasing order of mixture parameter w , and in Tables 3–5 we give the rank for each component that such an ordering would yield. However, we have found it instructive to divide θ_{134} 's components into three groups, and to manually reorder the components of

each group, in order to elucidate several prominent features of the probability landscape the DM represents. It may be possible to automate such a grouping and ordering using distance measures between Dirichlet distributions (Rauber et al., 2008). Developing such a method would be of interest, but for our present purposes it would provide only a distraction.

Table 4

Perhaps the most important feature of θ_{134} is represented by the 98 components of Group A (Table 3). As one moves from one component to the next within this group, the center of mass usually changes only slightly, and in a relatively continuous manner. The superposition of the probability hills represented by individual components can thus be visualized as a probability ridge threading its way through Ω_{20} , with several minor spurs.

The second feature of θ_{134} is represented by the 17 components of Group B (Table 4), which share two main properties: a preference for the amino acids glycine (G) and/or proline (P), and an aversion to hydrophobic amino acids. This group of components can be seen as a secondary ridge, separated from the first.

Table 5

A third feature is represented by the 19 components of Group C (Table 5). These components strongly favor a single amino acid, without clear secondary preferences that would attach them to the major ridge of Group A. Of note is the last component in this group, whose concentration parameter α is very close to 0. This implies a probability density concentrated almost completely near the vertices of Ω_{20} . It is a peculiarity of DMs that such densities can be approximated either by a single component with small α and probability mass dominated by several letters, or by the superposition of multiple components each with large α and probability mass dominated by a single letter. Thus this last component can be seen, in essence, as a formal alternative to the type of probability density represented by the other components of Group C.

In general, it may seem surprising that the data will support the identification of the 134 Dirichlet components shown in Tables 3–5, not to mention the > 600 components of many DMs with greater Δ . However, the $> 300,000$ columns in S_{UCSC} can associate on average > 500 columns to each of 600 components, and this much data is able to support relatively fine distinctions between similar probability densities.

When one compares different mixtures returned by the DP-sampler, the overall shape of the probability densities they describe can be recognized as remarkably similar. In contrast, the parameters of the individual components that go into approximating this shape have no particular stability. For example, a point that is halfway between the crests of two components in one mixture may very well be at the crest of an individual component in another.

4 Conclusion

When a set of data is believed to be well described by a mixture distribution, but with an unknown number of components, the Dirichlet process may be applied to infer the mixture (Blei and Jordan, 2005). Because homologous positions within protein families have been fruitfully described by Dirichlet mixtures (Brown *et al.*, 1993; Sjölander *et al.*, 1996; Altschul *et al.*, 2010), we have sought here to infer such mixtures from multiple alignment data using a Gibbs sampling algorithm based upon the Dirichlet process. This required us to develop several technical innovations, because the Dirichlet process has not previously been applied to DMs.

In contrast to previous approaches (Brown *et al.*, 1993; Sjölander *et al.*, 1996; Ye *et al.*, 2011b), our DP-sampler yields many hundreds of Dirichlet components when applied to real multiple alignment data. To understand these results, one should recognize that DMs are employed to model proteins primarily for mathematical as opposed to biological reasons: with Bayesian analysis, the posterior of a DM prior is still a DM (Brown *et al.*, 1993; Sjölander *et al.*, 1996; Altschul *et al.*, 2010). The DM formalism suggests the metaphor of discrete probability hills in Ω_{20} , each representing a category for classifying protein positions. However, the actual probability topography in Ω_{20} that describes proteins appears to be qualitatively different, having, for example, long probability ridges. To model such features well using Dirichlet components requires a large number of them, with closely spaced centers of mass. Our analysis suggests there is no “correct” number of components or categories for describing the probability distribution over Ω_{20} implied by proteins. Instead, when the MDL principle is applied, steadily increasing amounts of data should support steadily increasing numbers of components. However, as the number of components grows, there is also steadily

diminishing improvement, as measured by Δ , in modeling the underlying probability distribution.

The DP-sampler is able to find DMs that model multiple alignment data better than do those mixtures found by previously proposed methods. A key to its relative success is its ability to seed new components with columns that are not modeled well by any existing components, but to abandon components that do not then attract other columns. This fosters a much more efficient search of the very high-dimensional Dirichlet-mixture space than does seeding the space with random starting positions.

Although existing multiple alignment data sets may support DMs with over 500 components, speed considerations may favor smaller mixtures for use in practical sequence comparison algorithms. The DP-sampler can generate mixtures of many different sizes, to facilitate such a tradeoff.

At a deeper level, the DP-sampler provides a new perspective on the topography of protein space. This perspective suggests that the amino acid preferences at individual protein positions should, in general, be thought of not as falling into one of several categories, but rather as arrayed along a continuum. These preferences, represented by points in Ω_{20} , fall mainly near a long, almost one-dimensional probability ridge winding through the space. This perspective may suggest interesting questions for further investigation. For example, multiple alignment columns that imply similar high likelihoods for components situated far from one another along the ridge might imply either misalignment, or the presence of distinct protein subfamilies within the alignment.

Acknowledgments

Jordan Boyd-Graber is supported by NSF grant #1018625. Stephen Altschul is supported by the Intramural Research Program of the National Library of Medicine at the National Institutes of Health.

References

- Aldous, D. 1985. Exchangeability and related topics. In Ecole d’Ete de Probabilities de Saint-Flour XIII 1983, 1–198. Springer.
- Altschul, S. F., Wootton, J. C., Zaslavsky, E. et al. 2010. The construction and use of log-odds substitution scores for multiple sequence alignment. PLoS Comp. Biol. 6, e1000852.
- Antoniak, C. E. 1974. Mixtures of Dirichlet processes with applications to Bayesian nonparametric problems. Ann. Stat. 2, 1152–1174.
- Beal, M. J., Ghahramani, Z., and Rasmussen, C. E. 2002. The infinite hidden Markov model. In Adv. Neural Inf. Process. Syst. MIT Press.
- Blei, D. M. and Jordan, M. I. 2005. Variational inference for Dirichlet process mixtures. J. Bayesian Anal. 1, 121–144.
- Blei, D. M., Ng, A., and Jordan, M. 2003. Latent Dirichlet allocation. J. Mach. Learn. Res. 3, 993–1022.
- Brown, M., Hughey, R., Krogh, A. et al. 1993. Using Dirichlet mixture priors to derive hidden Markov models for protein families. In Proc. First International Conf. on Intell. Syst. for Mol. Biol., 47–55. AAAI Press.
- Edgar, R. C. and Sjölander, K. 2004. A comparison of scoring functions for protein sequence profile alignment. Bioinformatics 20, 1301–1308.
- Engen, S. 1975. A note on the geometric series as a species frequency model. Biometrika 62, pp. 697–699.

- Ferguson, T. S. 1973. A Bayesian analysis of some nonparametric problems. Ann. Stat. 1, 209–230.
- Gershman, S. J. and Blei, D. M. 2012. A tutorial on Bayesian nonparametric models. J. Math. Psych. 56, 1–12.
- Griffiths, R. 1980. Lines of descent in the diffusion approximation of neutral Wright-Fisher models. Theoretical Population Biol. 17, 37–50.
- Grünwald, P. D. 2007. The Minimum Description Length Principle, MIT Press Books, volume 1. The MIT Press.
- Hannah, L., Blei, D. M., and Powell, W. B. 2011. Dirichlet process mixtures of generalized linear models. J. Mach. Learn. Res. 12, 1923–1953.
- Hardisty, E., Boyd-Graber, J., and Resnik, P. 2010. Modeling perspective using adaptor grammars. In In Proc. Conf. Empirical Methods in Natural Language Processing.
- Lewis, D. D. 1998. Naive (Bayes) at forty: The independence assumption in information retrieval. In European Conf. on Mach. Learn., ECML '98, 4–15.
- McCloskey, J. 1965. A Model for the Distribution of Individuals by Species in an Environment. Ph.D. thesis, Department of Statistics, Michigan State University. (Unpublished).
- Minka, T. P. 2000. Estimating a Dirichlet distribution. Technical report, Microsoft.
- Müller, P. and Quintana, F. A. 2004. Nonparametric Bayesian data analysis. Statist. Sci. 19, 95–110.
- Neal, R. M. 2003. Slice sampling. Ann. Stat. 31, 705–767.

- Pitman, J. and Yor, M. 1997. The two-parameter Poisson-Dirichlet distribution derived from a stable subordinator. Ann. Probab. 25, 855–900.
- Rauber, T. W., Braun, T., and Berns, K. 2008. Probabilistic distance measures of the Dirichlet and Beta distributions. Pattern Recogn. 41, 637–645.
- Sethuraman, J. 1994. A constructive definition of Dirichlet priors. Statistica Sinica 4, 639–650.
- Sjölander, K., Karplus, K., Brown, M. et al. 1996. Dirichlet mixtures: a method for improved detection of weak but significant protein sequence homology. Comput. Appl. Biosci. 12, 327–345.
- Swanson, R. 1984. A vector representation for amino acid sequences. Bull. Math. Biol. 46, 623–639.
- Teh, Y. W., Jordan, M. I., Beal, M. J. et al. 2006. Hierarchical Dirichlet processes. J. Amer. Statist. Assoc. 101, 1566–1581.
- Wallach, H. M. 2008. Structured Topic Models for Language. Ph.D. thesis, University of Cambridge.
- Ye, X., Wang, G., and Altschul, S. F. 2011a. An assessment of substitution scores for protein profile-profile comparison. Bioinformatics 27, 3356–3363.
- Ye, X., Yu, Y.-K., and Altschul, S. F. 2011b. On the inference of Dirichlet mixture priors for protein sequence comparison. J. Comput. Biol. 18, 941–954.
- Yu, Y.-K. and Altschul, S. F. 2011. The complexity of the Dirichlet model for multiple alignment data. J. Comput. Biol. 18, 925–939.

List of Figures

- 1 Density plots for four Dirichlet distributions. The densities are over the triangular simplex that represents multinomial distributions over three letters, and demonstrate how different Dirichlet components can accommodate variable concentrations. Darker coloring denotes higher probability density. (a) Dirichlet parameters that are all 1.0 yield a uniform density over multinomial distributions. (b) Dirichlet parameters that are all greater than 1.0 yield a density concentrated near the mean \vec{q} , in this case (0.6250, 0.0625, 0.3125). (c) and (d) Dirichlet parameters that are all less than 1.0 yield a density concentrated near the edges and corners of the simplex. Such a density favors sparse multinomials, in which only a subset of letters has appreciable probability. Symmetric (c) and asymmetric (d) cases are shown. 37
- 2 Graphical model representing the Dirichlet process Dirichlet mixture model proposed in this paper. Nodes represent variables; shaded nodes are observed; edges show probabilistic dependencies; and plates denote replication. 38
- 3 Δ s for Dirichlet mixtures found by the DP-sampler. Using parameters $\beta = 400$, and $\gamma = 100$, the DP-sampler was run on S_{UCSC} for 1000 iterations. After every 10th iteration, the MDL principle was applied to the components returned, to find a DM with optimal Δ . Black crosses indicate Δ s that are greater than those for all previous iterations; red circles, others. For iteration 10, $\Delta = 1.0666$ is off the scale. 39

- 4 Number of Dirichlet mixture components returned. The DP-sampler was run as described in the caption to Figure 3. Red squares show the number of components returned by the DP-sampler after every 10th iteration. Black circles show the number of components in DMs supported by the MDL principle; their corresponding Δ s are shown in Figure 3. . 40
- 5 Best Δ for a specified number of mixture components. The DP-sampler was run as described in the caption to Figure 3. Shown are the best Δ s found, during any iteration, for DMs with a given number of components. 41




Figure 1: Density plots for four Dirichlet distributions. The densities are over the triangular simplex that represents multinomial distributions over three letters, and demonstrate how different Dirichlet components can accommodate variable concentrations. Darker coloring denotes higher probability density. (a) Dirichlet parameters that are all 1.0 yield a uniform density over multinomial distributions. (b) Dirichlet parameters that are all greater than 1.0 yield a density concentrated near the mean \vec{q} , in this case $(0.6250, 0.0625, 0.3125)$. (c) and (d) Dirichlet parameters that are all less than 1.0 yield a density concentrated near the edges and corners of the simplex. Such a density favors sparse multinomials, in which only a subset of letters has appreciable probability. Symmetric (c) and asymmetric (d) cases are shown.

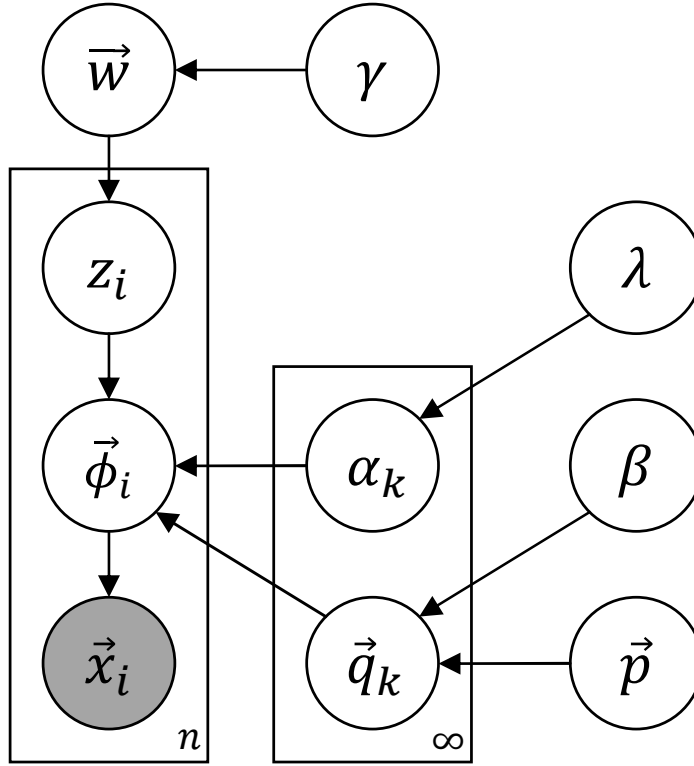


Figure 2: Graphical model representing the Dirichlet process Dirichlet mixture model proposed in this paper. Nodes represent variables; shaded nodes are observed; edges show probabilistic dependencies; and plates denote replication.




Figure 3: Δ s for Dirichlet mixtures found by the DP-sampler. Using parameters $\beta = 400$, and $\gamma = 100$, the DP-sampler was run on S_{UCSC} for 1000 iterations. After every 10th iteration, the MDL principle was applied to the components returned, to find a DM with optimal Δ . Black crosses indicate Δ s that are greater than those for all previous iterations; red circles, others. For iteration 10, $\Delta = 1.0666$ is off the scale.




Figure 4: Number of Dirichlet mixture components returned. The DP-sampler was run as described in the caption to Figure 3. Red squares show the number of components returned by the DP-sampler after every 10th iteration. Black circles show the number of components in DMs supported by the MDL principle; their corresponding Δ s are shown in Figure 3.




Figure 5: Best Δ for a specified number of mixture components. The DP-sampler was run as described in the caption to Figure 3. Shown are the best Δ s found, during any iteration, for DMs with a given number of components.

List of Tables

1	Effect of Dirichlet Process Parameters	43
2	The DP-Sampler, with Slice Sampling for γ	44
3	Parameters of θ_{134} : Group A	45
4	Parameters of θ_{134} : Group B	46
5	Parameters of θ_{134} : Group C	47

Table 1: Effect of Dirichlet Process Parameters

β	γ	Best Δ (bits/a.a.)	Number of components	Iteration found	β	γ	Best Δ (bits/a.a.)	Number of components	Iteration found
100	5	1.0756	470	950	600	5	1.0757	361	980
	10	1.0758	520	860		10	1.0760	406	960
	20	1.0759	644	980		20	1.0762	471	930
	40	1.0760	689	980		40	1.0762	481	980
	60	1.0760	720	800		60	1.0763	533	910
	80	1.0760	717	630		80	1.0762	526	900
	100	1.0760	808	1000		100	1.0762	541	780
200	5	1.0757	449	1000	800	5	1.0757	341	1000
	10	1.0759	498	840		10	1.0759	378	1000
	20	1.0761	586	960		20	1.0761	431	960
	40	1.0761	597	600		40	1.0761	466	830
	60	1.0762	665	750		60	1.0762	472	830
	80	1.0762	709	870		80	1.0762	471	730
	100	1.0762	679	590		100	1.0762	499	760
400	5	1.0757	400	980	1000	5	1.0755	314	1000
	10	1.0760	452	910		10	1.0758	350	1000
	20	1.0762	505	960		20	1.0759	375	990
	40	1.0763	562	860		40	1.0760	429	860
	60	1.0763	588	1000		60	1.0761	433	910
	80	1.0763	603	990		80	1.0761	447	910
	100	1.0763	623	940		100	1.0761	444	990

Table 2: The DP-Sampler, with Slice Sampling for γ

β	Mean and standard deviation of γ		Best Δ (bits/a.a.)	Number of components	Iteration found	γ
100	199.7 \pm	13.0	1.0760	767	280	200
200	183.9 \pm	6.8	1.0762	721	580	166
400	129.6 \pm	4.3	1.0763	608	790	128
600	95.5 \pm	3.9	1.0762	537	930	94
800	82.2 \pm	3.5	1.0762	482	990	83
1000	66.8 \pm	2.8	1.0760	442	940	69

Table 3: Parameters of θ_{134} : Group A

Rank	w (%)	α_k	RKQEDNHWYFMLIVCTSAGP	Rank	w (%)	α_k	RKQEDNHWYFMLIVCTSAGP
69	0.51	30.7	R .-=- -	93	0.40	24.9	----- FmL . -
23	1.20	26.7	R+	65	0.55	52.7	----- fml+ .-.-
124	0.26	35.3	K.-.-.-.-	92	0.40	34.0	----- fmliv+ . .
15	1.49	27.0	rK+ . -.- .-.-	4	2.56	37.2-.. +mL+ .. -
3	2.82	27.0	rk+ - + -	64	0.57	32.6	-----+mLI -
89	0.41	0.4	RKq - +- -	11	1.67	49.0	.-.-.-. ++liv . -
24	1.16	33.0	+++ -.-.-.- +a .	125	0.25	6.6	-- -.-.-. M+ .-
7	1.91	62.7	rkq+ ..-	43	0.76	14.8	.- -=-. ML+ + . -
2	3.18	59.5	++++ .	39	0.82	6.6	----- mL+ .-
91	0.41	164.5	+kqe+ .----- a..	16	1.44	67.9	++++ .
6	1.95	106.3	+kqe+ ..-.-.-. .	76	0.48	22.2	----- mliv - =
18	1.37	37.2	+kqE+ - - - - . .	105	0.32	28.0	----- mli++ ,a.=
25	1.13	36.1	+k+ +n - - - - +	35	0.97	61.8	----- +L .-
19	1.33	97.6	+++++ +	54	0.65	82.9	----- +lIv --.=
41	0.80	74.4	++edn -.-.-.-	99	0.37	47.9	----- lIV.-.=.
60	0.61	22.7	Q+	29	1.00	22.3	----- +Iv.-=-
83	0.45	6.9	. qE+ .- - - .	106	0.32	3.5	----- IV--=-
51	0.67	57.6	. qEd - - - - . -	72	0.49	54.4	----- IV -=-
5	2.15	34.3	+E	42	0.78	52.4	----- IV .-.-
85	0.44	43.2	-- E .- - - - -	8	1.86	10.4	.-... iv . -
95	0.39	63.2	+e+ - - - - s -	9	1.85	37.2	.. iv ..
27	1.04	107.4	+Ed - - - - .	71	0.50	70.9	----- iV .- -
101	0.35	0.4	=- ED = - - - =	46	0.71	17.4	----- iV - =
86	0.44	43.3	eD - - - - -	61	0.59	36.3	----- iV+ ,a.-
129	0.21	23.0	-- eD - - - - -	22	1.22	23.4	.-.-. + vT+ ..
10	1.68	38.4	+Dn	31	0.99	4.7	-=.-=... m +C a .
126	0.24	13.2	-- . D - -.-.-. ++ .	34	0.97	34.7	----- ++ +c a -
79	0.47	61.8	Dn - - - - . + .	68	0.52	34.9	----- . +c A.-
117	0.29	24.9	-- DN - - - - . -	32	0.98	34.9	.-.-. + A -
48	0.68	26.8	dN+-.-.-.-. .	74	0.48	9.7	=====.-. . vCTsa..
109	0.32	25.3	---- N = - - - .-.-	73	0.48	38.1	.-... c+sa .
98	0.37	29.9	.. -.N+. + .-	131	0.19	22.4	-----.------c+Sa -
17	1.38	27.8	+++ nh y .- . .	103	0.34	5.2	-----.------c sA+.
63	0.58	70.7	++++ . . +	90	0.41	0.4	- -. - - - - -C+s g+
70	0.51	21.5	. . . H y	21	1.28	13.6	.. . ++ + ...++s
58	0.62	4.7	hWYf --. . .	102	0.35	13.1	-----.------T+---
96	0.38	1.4	-----+WYF-----	47	0.69	27.3	.-.-. -.-.-.-. Ts ..
13	1.63	23.8	...-.+wYF . . -	97	0.38	35.6	.. +n.------Ts. .
118	0.29	27.9	-----W+ .- - - -	44	0.75	2.7	- nh--= -- ts +
77	0.47	26.6	Wy+ ..	12	1.67	44.1	++ts.-.
130	0.19	38.5	----- WyF . . . -	28	1.03	49.4	n +s
114	0.30	24.8	----- wYF . . -	94	0.39	20.3	----- +S.-
1	3.44	29.6	. wyf++ .	75	0.48	23.7	. - - - - - +S
128	0.21	21.0	-----W+fm++ .-.-	116	0.29	11.0	--.. - - - - - s G
80	0.47	32.6	----- +Y+-----	120	0.28	46.1	--.-. -.-.-.-. saG.
38	0.84	24.6	-----+yF . . -	132	0.18	39.3	-----.-. +. AG-
81	0.46	11.3	----- +Yf++iv --.-	112	0.31	24.2	-----.- - aG-
123	0.27	11.7	. + y+m .	121	0.27	90.2	-----.-.- G-
53	0.66	33.1	----- +F++i+ .-.-	115	0.29	14.6	--.-.-.-.-.- a P

Table 4: Parameters of θ_{134} : Group B

Rank	w (%)	α	RKQEDNHWYFMLIVCTSAGP									
100	0.36	32.5	+k+	n	-----	.	.	G.				
82	0.46	38.1	.	dn	-.-----	.	.	G.				
78	0.47	100.0	.	n	=====	-	G.					
55	0.63	83.2		++	-----		G					
30	1.00	50.3		+	...-. .		G					
57	0.62	82.6			.-----	.	.	G				
113	0.31	43.1			-----		gP					
45	0.72	75.9		+d+	..-----	.	+	+p				
108	0.32	31.7	.	d+	-----	.	s	P				
127	0.21	77.4		d+	-----	.	ts.	p				
56	0.63	69.9		ed	-----		P					
110	0.31	84.8	+k+e+		-----		p					
119	0.28	9.2	rk	d+	==	==	.	p				
50	0.67	41.6	rk+		-.		p					
33	0.98	85.6		+		p					
59	0.62	66.7	+	+	-----		.P					
87	0.44	48.5	.		.-----	.	.	P				

Table 5: Parameters of θ_{134} : Group C

Rank	w (%)	α	RKQEDNHWFMLIVCTSAGP
111	0.31	16.3	Q -. --- .----..----
67	0.52	52.8	.. D -----.. ...
88	0.41	60.4	==-.D.-=====
122	0.27	34.9	--.---H-.-----
133	0.18	24.5C ...
134	0.16	59.0	==-=-----C-----
14	1.60	41.8	. a .
66	0.55	40.3	=====-----.. . A -
26	1.06	43.8	g
62	0.59	27.9	.-.---.. G.
49	0.68	112.4	-.-. . =-----.-G-
36	0.94	80.3	=====-----G=
20	1.32	66.2	p
40	0.82	44.8	.. .P
37	0.93	42.9 P
107	0.32	17.3	-----.. . -P
84	0.44	62.5-..... .. .P
52	0.66	51.7	-----.-P
104	0.34	0.0	. --H-. ===C-. -Gp