



Probability Distributions: Continuous

Introduction to Data Science Algorithms

Jordan Boyd-Graber and Michael Paul

SEPTEMBER 15, 2016

Combining Discrete and Continuous Distributions

- We can chain together two distributions
- E.g., imagine your multinomial distribution came from a Dirichlet
- Often called “Bayesian Data Analysis”
- This why explain why “add one” Laplace smoothing isn’t crazy

Why Bayesian

- Imagine you have vector of counts \vec{n} that come from multinomial $\vec{\theta}$. This multinomial comes from a Dirichlet with parameter $\vec{\alpha}$. (Chain rule)

$$p(\vec{n}) = p(\vec{n} | \theta) p(\vec{\theta} | \vec{\alpha}) \quad (1)$$

- Now let's assume that you see some counts \vec{n} . You want to know what the multinomial distribution parameter looks like.

$$p(\vec{\theta} | \vec{n}, \vec{\alpha}) \quad (2)$$

Conjugacy

- If $\vec{\theta} \sim \text{Dir}(\cdot|\alpha)$, $\vec{w} \sim \text{Mult}(\cdot|\theta)$, and $n_k = |\{w_i : w_i = k\}|$ then

$$p(\theta|\alpha, \vec{w}) \propto p(\vec{w}|\theta)p(\theta|\alpha) \quad (3)$$

$$\propto \prod_k \theta^{n_k} \prod_k \theta^{\alpha_k - 1} \quad (4)$$

$$\propto \prod_k \theta^{\alpha_k + n_k - 1} \quad (5)$$

- Conjugacy: this **posterior** has the same form as the **prior**
- In fact, it looks like you're just adding counts!

Conjugacy

- If $\vec{\theta} \sim \text{Dir}(\alpha)$, $\vec{w} \sim \text{Mult}(\theta)$, and $n_k = |\{w_i : w_i = k\}|$ then

$$p(\theta|\alpha, \vec{w}) \propto p(\vec{w}|\theta)p(\theta|\alpha) \quad (3)$$

$$\propto \prod_k \theta^{n_k} \prod_k \theta^{\alpha_k - 1} \quad (4)$$

$$\propto \prod_k \theta^{\alpha_k + n_k - 1} \quad (5)$$

- Conjugacy: this **posterior** has the same form as the **prior**
- In fact, it looks like you're just adding counts!

Why add one?

- The count that we add is equivalent to the Dirichlet parameter
- What does this mean in the case of Dirichlet distribution?

$$f(\theta) = \frac{\prod_{k=1}^K \Gamma(\alpha_k)}{\Gamma(\sum_{k=1}^K \alpha_k)} \prod_{k=1}^K \theta_k^{\alpha_k - 1}$$

Why add one?

- The count that we add is equivalent to the Dirichlet parameter
- What does this mean in the case of Dirichlet distribution?

$$f(\theta) = \frac{\prod_{k=1}^K \Gamma(\alpha_k)}{\Gamma(\sum_{k=1}^K \alpha_k)} \prod_{k=1}^K \theta_k^{\alpha_k - 1}$$

- Uniform distribution! Doesn't matter what x is.

Next time

- Drawing from and plotting various distributions
- Be sure to bring laptops