



Introduction

Natural Language Processing: Jordan
Boyd-Graber

University of Colorado Boulder

AUGUST 25, 2014

Roadmap

By the end of this class you should ...

- Be able to give examples of where NLP is used
- Understand the workings of the course
- Know me and each other a little better

Outline

Natural Language Processing

Administrivia and Introductions

Introducing Python and NLTK

Where do you see natural language processing?

Machine Learning is Doing Great!



- Can drive a million miles without an accident
- Can beat any living chess player



Machine Learning is Doing Great!



- Can drive a million miles without an accident
- Can beat any living chess player
- Automated call center vs. five-year old?



Machine Learning is Doing Great!



- Can drive a million miles without an accident
- Can beat any living chess player
- Automated call center vs. five-year old?
- We'll learn why we're so far away

What's Natural Language Processing

- Computational approaches to understand, generate, and process natural language
- Cross-discipline
 - Computer science: implement algorithms
 - Linguistics: develop theory / data
 - Statistics: learn patterns from data
 - Experts in specific languages: get a computer to handle a new language
 - Psychologists: how does our brain process language
 - Sociologists: how do social constraints change how we process language

What computational linguistics can do!

Automatic solutions to . . .

- Explain why the “ly” in “**ally**” and “quick**ly**” are different (morphology)
- Tell the difference in category between “**water** the flowers” and “drink the **water**” (part of speech tagging)
- Why “saw the sun with the telescope” is different from “saw the astronomer with the telescope”
- Translate “My hovercraft is full of eels” into Hungarian (machine translation)

Outline

Natural Language Processing

Administrivia and Introductions

Introducing Python and NLTK


What you need for this course

- Interactive classroom: come ready with questions and participate in exercises
- Helps to have a laptop to bring to class
- Math background
 - Will ask you to manipulate equations
 - Will expect you to be able to do basic derivatives
 - Work with functions like exponentiation and logs
 - Probability: review next week (hugely important)
- Computer / programming skills
 - You will need to write python programs
 - You will need to interact with a Unix command line
 - You will need to interact with data files

Administrivia

- Sign up on Piazza (use a photo)
 - Another way to earn participation
 - CAETE: only way to earn participation
- Keep track of course webpage
http://www.cs.colorado.edu/~jbg/teaching/CSCI_5832/
- Homeworks (still working on handin mechanism)
- 5 late days
- Let me know about special needs

Course reading



- Colorado roots: written here
- We will provide reading assignments, mostly from the book. (Read them **before** associated class.)
- The reading will cover more than we cover in class.
- Don't buy the first edition

Communicating with Piazza

We will use Piazza to manage all communication

`https://piazza.com/umd/fall2013/cmsc723/home`

- Questions answered within 1 day (hopefully sooner)
- Hosts discussions among yourselves
- Use for any kind of technical question
- Use for **most** administrative questions
- Can use to send us private questions too

How to ask for help

- Explain what you're trying to do
- Give a minimal example
 - Someone else should be able to replicate the problem easily
 - Shouldn't require any data / information that only you have
- Explain what you **think** should happen
- Explain what you get instead (copy / paste or screenshot if you can)
- Explain what else you've tried

Me

- First year assistant professor
 - Office: 111B ECCS
- Was formerly a professor at University of Maryland
- Research: topic models, question answering, machine translation
- Third time teaching the class (or something similar)
- Born in Colorado (where all my family live)
- Grew up in Iowa (hometown: Keokuk, Iowa)
- Went to high school in Arkansas
- Undergrad in California
- Grad school in New Jersey
- Brief jobs in between:
 - Working on electronic dictionary in Berlin
 - Worked on Google Books in New York
- ying / jbg / jordan / boyd-graber

Grader

- Fortunate enough to have a grader, Skatje Myers
- **Do** address grading questions to her
- **Don't** seek her out if you're confused about a concept
- Instead, schedule a 1-on-1 appointment with me or come to office hours

Outline


Natural Language Processing

Administrivia and Introductions

Introducing Python and NLTK

Why Python?

- Easy to learn
- Widespread
- Can be fast if you need it (cython)



Why NLTK?

- Handy code for accessing data
- Implementations of standard algorithms
- Easy to quickly process text and try things out

Why NLTK?

- Handy code for accessing data
- Implementations of standard algorithms
- Easy to quickly process text and try things out
- Chapter 1 of NLTK book
- Ask questions on Piazza

Next time ...

- Install NLTK on your computer
- Come to class ready to work through some Python examples
- Also look at homework 1 and see if you have any questions